Logistic Regression & Neural Networks

Saeed Saremi

Assigned reading: 5.3.1, 5.4.{1, 2, 3, 4}, 6.{1, 2}

September 19, 2024

a summary of the previous lecture

- We "separated" the problem of classification, of learning the mapping $f : \mathcal{X} \to \{0, \ldots, K-1\}$, into an inference step and the decision step.[1]
- We first took a generative viewpoint, assuming $X|k \sim \mathcal{N}(\mu_k, \Sigma)$ and found out for two-class classification the posterior $p(k|x)$ takes the form of the logistic function:

$$p(k|x) = \frac{1}{1 + e^{-z}},$$

where z is linear function:

$$z = \theta^\top x + \theta_0,$$

where θ is depends on $\{\mu_k\}_{k=0}^{K-1}$ and Σ.

- In the generative approach μ_k and Σ are estimated from the data.
- We learnt than one can parametrize the posterior distribution directly and by some considerations, logistic function again appears. In this discriminative approach, the parameters (θ, θ_0) should be estimated directly with maximum likelihood.

[1] For simple criterion of minimizing the probability of misclassification, we arrive at

$$k = \underset{k}{\operatorname{argmax}}\, p(k|x).$$

outline

- on the geometrical meaning of θ in the logistic regression
- multi-class ($K > 2$) extension of the logistic function
- the cross-entropy loss for discriminative models and its gradient
  - neural networks

The GEOMETRY of θ in LOGISTIC regression

$$p_\theta(1|x) = \frac{1}{1+\exp(-\theta^T x - \theta_0)} \quad (1)$$

$$\left.\begin{aligned} x &= x_{\parallel} + x_{\perp} \\ \theta^T x_{\perp} &= 0 \end{aligned}\right\} \Rightarrow \text{adding } \alpha x_{\perp} \text{ to } x \text{ does not change } p_\theta(1|x)$$

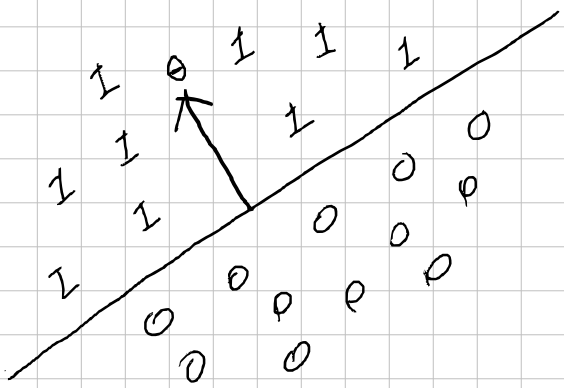

- θ is perpendicular to the decision boundary.
- with parametrization given in Eq. (1), θ "points" to class "1".
- θ_0 shifts the decision boundary.

MULTI-CLASS ($K > 2$) CLASSIFICATION

As before we start with the generative approach:

$$X \mid k \sim N(\mu_k, \Sigma)$$

$X \in \mathbb{R}^d$, $k \in \{1, \dots, K\}$, $\mu_k \in \mathbb{R}^d$, $\Sigma \in \mathbb{R}^{d \times d}$

$$p(k \mid x) = \frac{p(x \mid k)\, p(k)}{\sum_{j=1}^{K} p(x \mid j)\, p(j)}$$

$$= \frac{e^{a_k}}{\sum_{j=1}^{K} e^{a_j}}$$

$$p(x \mid k) = \frac{1}{(2\pi)^{d/2} |\Sigma|} e^{-\frac{1}{2}(x-\mu_k)^T \Sigma^{-1} (x-\mu_k)}$$

$$a_k = \underbrace{(\Sigma^{-1}\mu_k)^T}_{\theta_k} x \underbrace{- \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log p(k)}_{\theta_{0,k}}$$

$$a_k = \theta_k^T x + \theta_{0,k}$$

$$(a_1 \cdots a_K) \mapsto \left(\frac{e^{a_1}}{\sum e^{a_j}}, \dots, \frac{e^{a_K}}{\sum e^{a_j}} \right)$$

SOFTMAX FUNCTION

If $a_k \gg a_j$, $j \neq k$: $(a_1 \cdots a_K) \mapsto (00 \cdots 1 \cdots 0)$ (index k)

LEARNING in DISCRIMINATIVE models

We start with our linear two-class logistic regression model:

$$p_\theta(1|x) = \frac{1}{1 + \exp(-\theta^T x - \theta_0)}$$

The log-likelihood of this model for any $y \in \{0,1\}$ is given by:

$$\log\left(p_z^{\,y}(1-p_z)^{1-y}\right),$$

where z depends on θ & crucially p only depends on z.

We define the loss (something we minimize) as the negative log likelihood

$$\mathcal{L}(\theta) = -y \log p_z - (1-y)\log(1-p_z)$$

We need to know the gradient $\nabla_\theta \mathcal{L}(\theta)$ to optimize for parameters θ:

We use the chain rule:

$$\nabla_\theta \mathcal{L}(\theta) = \partial_z\left(-y \log p_z - (1-y)\log(1-p_z)\right) \nabla_\theta z$$

Second, we use the following

$$\partial_z p_z = \partial_z \left(\frac{1}{1+e^{-z}} \right) = \frac{e^{-z}}{(1+e^{-z})^2} = p_z (1 - p_z)$$

We therefore have:

$$\nabla_\theta \mathcal{L}(\theta) = \left(-y \frac{p_z(1-p_z)}{p_z} + (1-y) \frac{p_z(1-p_z)}{1-p_z} \right) \nabla_\theta z$$

$$= (p_z - y) \nabla_\theta z \longrightarrow x \quad \text{(for linear model)}$$

The final answer has a nice interpretation:

$$\theta \leftarrow \theta - \epsilon (p_z - y) x$$

$$\begin{cases} \theta \leftarrow \theta - \tilde{\epsilon}\, x & y = 0 \\ \theta \leftarrow \theta + \tilde{\epsilon}\, x & y = 1 \end{cases}$$

(This is expected from our geometrical picture of the logistic regression.)

θ is "pushed" in the direction of x for $y = 1$.

The reverse happens for $y = 0$.

NEURAL NETWORKS

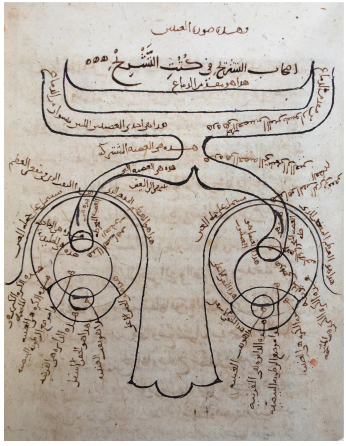

Figure: The oldest known drawing of the nervous system by Ibn al-Haytham (published in 1083)

Figure: Olfactory bulb, Camillo Golgi, 1875

Santiago Ramón Y Cajal & the neuron doctrine

Neuron as the discrete distinct entities in the brain as opposed to a continuous network.

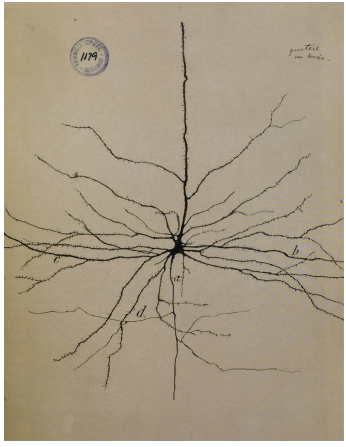

Figure: *pyramidal neuron.* Cajal, 1899

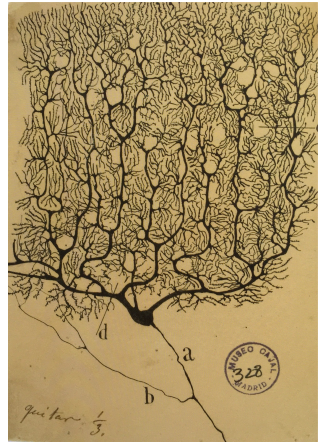

Figure: *purkinje neuron.* Cajal, 1899

The Nobel Prize in Physiology or Medicine 1906

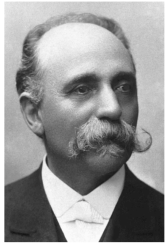

Photo from the Nobel Foundation archive.

Camillo Golgi

Prize share: 1/2

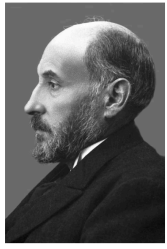

Photo from the Nobel Foundation archive.

Santiago Ramón y Cajal

Prize share: 1/2

The Nobel Prize in Physiology or Medicine 1906 was awarded jointly to Camillo Golgi and Santiago Ramón y Cajal "in recognition of their work on the structure of the nervous system"

network of neurons: axons, dendrites, and synapses

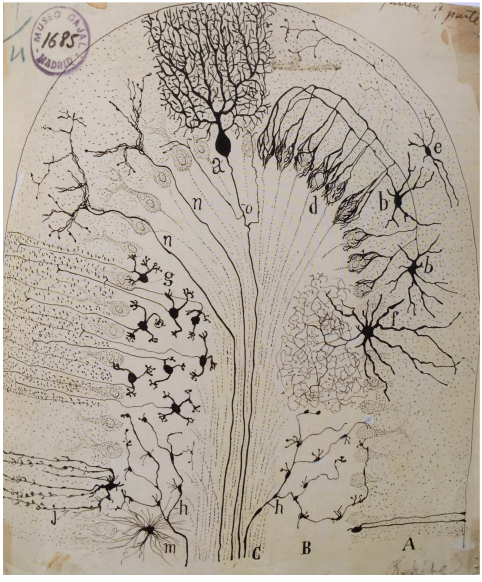

Figure: *network of neurons.* Cajal, 1899

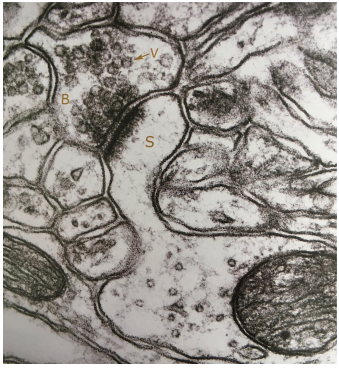

Figure: *Synapse.* Spacek and Harris, 2000

electricity in the brain

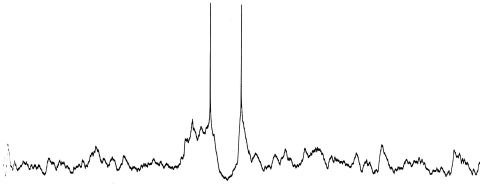

Figure: *Whole-cell recording in an awake rat.* Contantinople and Bruno, 2009.

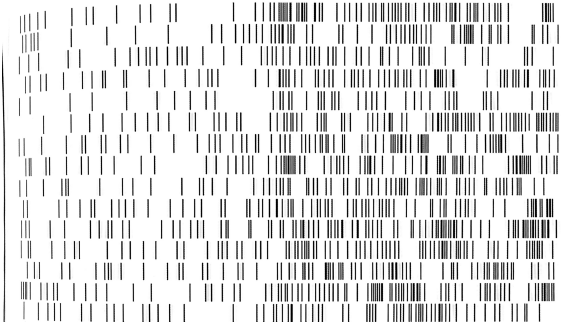

Figure: *Action potentials in a live monkey brain.* Saez and Salzman, 2009. (Each row, 4 seconds.)

biological neural networks v.s. deep neural networks

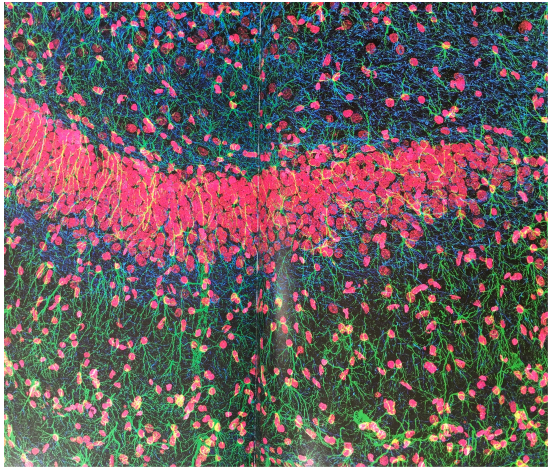

Figure: human brain: 10^{11} neurons, 10^{15} synapses

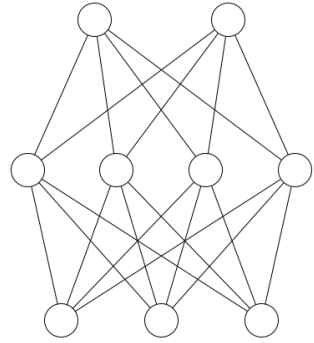

Figure: GPT-4: 10^6 neurons, 10^{11} parameters (weights), 100 layers

❓ What are we missing in this comparison?